Separate Anything You Describe

Xubo Liu, Qiuqiang Kong, Yan Zhao, Haohe Liu, Yi Yuan, Yuzhuo Liu,
Rui Xia, Yuxuan Wang, Mark D. Plumbley, Wenwu Wang
https://audio-agi.github.io/Separate-Anything-You-Describe

Abstract—Language-queried audio source separation (LASS) is a new paradigm for computational auditory scene analysis (CASA). LASS aims to separate a target sound from an audio mixture given a natural language query, which provides a natural and scalable interface for digital audio applications. Recent works on LASS, despite attaining promising separation performance on specific sources (e.g., musical instruments, limited classes of audio events), are unable to separate audio concepts in the open domain. In this work, we introduce AudioSep, a foundation model for open-domain audio source separation with natural language queries. We train AudioSep on large-scale multimodal datasets and extensively evaluate its capabilities on numerous tasks including audio event separation, musical instrument separation, and speech enhancement. AudioSep demonstrates strong separation performance and impressive zero-shot generalization ability using audio captions or text labels as queries, substantially outperforming previous audio-queried and language-queried sound separation models. For reproducibility of this work, we released the source code, evaluation benchmark and pre-trained model at: https://github.com/Audio-AGI/AudioSep.

Index Terms—sound separation, language-queried audio source separation (LASS), natural language processing.

I. INTRODUCTION

COMPUTATIONAL auditory scene analysis (CASA) [1] aims to design machine listening systems that perceive complex sound environments in a similar way to the human auditory system. As a fundamental research task for CASA, sound separation aims to separate real-world sound recordings into individual source tracks, also known as the "cocktail party problem" [2]. Sound separation has a wide range of applications, including audio event separation [3], [4], music source separation [5], and speech enhancement [6], [7].

Many previous works on sound separation mainly focus on separating one or a few sources, such as in speech enhancement [6], [7], speech separation [8], [9], and music source separation [5]. Recently, universal sound separation (USS) [4] has attracted a lot of research interest. USS aims to separate arbitrary sounds in real-world sound recordings. Separating every sound source from a mixture is challenging due to the wide variety of sound sources existing in the world. As an alternative, query-based sound separation (QSS) has been proposed which aims to separate specific sound sources conditioned on a piece of query information. QSS allows users

to extract desired audio sources which could be useful in many applications such as automatic audio editing [10] and multimedia content retrieval [11]. Using the query of different modalities such as vision [12]–[14], audio [15]–[18] or event labels [19]–[22] for sound separation has been investigated in the literature.

Recently, a new paradigm of QSS has been proposed, known as language-queried audio source separation (LASS) [3]. LASS is the task of separating arbitrary sound sources using natural language descriptions of the desired source. LASS provides a potentially useful tool for future digital audio applications, allowing users to extract desired audio sources via natural language instructions. The use of natural language queries offers significant advantages, compared to previous vision-queried [12]–[14] or audio-queried [15]–[18] methods, such as flexible and convenient acquisition of query information. Compared to label-queried [19]–[22] methods that usually pre-define a fixed set of label categories, LASS does not limit the scope of input queries and can be seamlessly generalized to open domain.

The challenge of learning a LASS system is associated with the complexity and variability of natural language expressions. The text query description could range from sophisticated descriptions of multiple sound sources, such as "*people speaking followed by music playing with a rhythmic beat*" to simple and compact phrases such as "*speech, music*". In addition, the same audio source can be delivered with diverse language expressions, such as "*music is being played with a rhythmic beat*" or "*an upbeat music melody is playing over and over again*". LASS not only requires these phrases and their relationships to be captured in the language description but also that one or more sound sources that match the language query should be separated from the audio mixture.

Our original approach [3] for LASS relies on supervised learning with labeled audio-text paired data. However, such annotated audio-text data is limited in size. To overcome the data scarcity issue, recent advancements have investigated training LASS with multimodal supervision [23]–[25]. The key idea behind this approach is to leverage multimodal contrastive pre-training models such as the contrastive language-image pretraining (CLIP) model, as the query encoder. As contrastive learning is capable of aligning text embedding with other modalities (e.g., vision), it enables the training of the LASS system using data-rich modalities and facilitates inference with text in a zero-shot mode. However, existing LASS methods leverage small-scale data for training and focus on separating restricted source types, such as musical instruments and a limited set of sound events. The potential to generalize LASS to open-domain scenarios, such as hundreds of real-world

Xubo Liu, Haohe Liu, Yi Yuan, Mark D. Plumbley, and Wenwu Wang are with the Centre for Vision, Speech and Signal Processing (CVSSP), University of Surrey, Guilford, UK. Email: {xubo.liu, haohe.liu, yi.yuan, m.plumbley, w.wang}@surrey.ac.uk.

Qiuqiang Kong, Yan Zhao, Yuzhuo Liu, Rui Xia, and Yuxuan Wang: are with the Speech, Audio & Music Intelligence (SAMI) Group, ByteDance Inc. Email: {kongqiuqiang, zhao.yan, liuyuzhuo.999, rui.xia, wangyuxuan.11}@bytedance.com.

sound sources, has not yet been fully explored. Furthermore, previous works on LASS evaluate model performance on each domain-specific test set, which makes future comparison and reproduction inconvenient and inconsistent.

In this work, our goal is to establish a foundation model for sound separation with natural language descriptions. Our focus is on the development of a pre-trained sound separation model, leveraging large-scale datasets, to enable robust generalization in open-domain scenarios. With this model, we aim to holistically address the separation of a diverse range of sound sources. This work constitutes an extension of our previous research, as presented in the Interspeech proceeding [3]. The contribution of this work includes:

- We introduce AudioSep, a foundation model for open-domain, universal sound separation with natural language queries. AudioSep is trained with large-scale audio datasets and has shown strong separation performance and impressive zero-shot generalization capabilities. LASS-Net model presented in the Interspeech work did not perform well in open-domain scenarios.

- We extensively evaluate AudioSep. We construct a comprehensive evaluation benchmark for LASS research, involving numerous sound separation tasks such as audio event separation, musical instrument separation, and speech enhancement. We show that AudioSep substantially outperforms off-the-shelf audio-queried sound separation and state-of-the-art LASS models.

- We conduct in-depth ablation studies to investigate the impact of scaling up AudioSep using large-scale multimodal supervision [23]–[25]. Our findings provide valuable insights for future research.

- We released the source code, evaluation benchmark, and pre-trained model at: https://github.com/Audio-AGI/AudioSep to promote research in this area.

II. RELATED WORK

A. Universal sound separation

Sound separation is a fundamental research task in CASA, which aims to separate real-world sound recordings into individual source tracks. A substantial amount of previous research has been concentrated on domain-specific sound separation, focusing primarily on areas such as speech [8], [9] or music [5]. Universal sound separation (USS) [4] aims to separate a mixture of arbitrary sound sources in terms of their classes. The challenge inherent in USS is the diversity of sound classes in real-world scenarios, which increases the difficulty of separating all of these sound sources with a single sound separation system. The work in [4] reported promising results on separating arbitrary sounds using permutation invariant training (PIT) [26], a supervised method initially designed for speech separation. The PIT method uses synthetic training mixtures simulated from single-source ground truth, performing sub-optimally due to a mismatch in the distribution between these synthetic mixtures and real-world sound recordings. Furthermore, it is not feasible to record a large database of single sources for PIT, as such sound recordings are often tainted by cross-talk. An unsupervised

method called mixture invariant training (MixIT) [27] was proposed for sound separation using noisy audio mixtures. MixIT has achieved competitive performance compared to supervised methods (e.g., PIT), showing substantial improvements in reverberant sound separation performance. Both PIT and MixIT methods need a post-selection process to classify separated sources into specific sound classes.

B. Query-based sound separation

Query-based sound separation (QSS), also known as target source extraction, aims to separate a specific source from an audio mixture given some query information. Existing QSS approaches could be divided into three categories: vision-queried, audio-queried, and label-queried. We will introduce each one below.

1) Vision-queried sound separation: In the computer vision community, there has been active research focusing on utilizing visual information as queries to extract target sounds in speech [28], [29], music [12], and acoustic events [13]. Audio-Scope [14] has been recently proposed to perform on-screen sound separation based on the MixIT method. Such vision-queried approaches are beneficial for automatically decomposing sound sources from audio-visual video data. However, their performance is affected by the dynamic visibility conditions of visual objects, and can be degraded in environments with visual occlusion or low-light conditions. In addition, video data often contains off-screen sounds. Learning from noisy audio-visual data is a key challenge for vision-queried sound separation systems [14], [23].

2) Audio-queried sound separation: Another line of research leverages the audio modality as a query to separate acoustically similar sounds. Recent studies [17], [18], [30] have addressed the problem of using one or a few examples of a target source as the query to separate a particular sound source: this is known as one-shot or few-shot sound separation. These methods separate the target sound conditioned on the average audio embedding of a few audio examples of the target source, which requires labeled single sources for calculating query embedding during training. The work in [15], [16] proposes to train the audio-queried sound separation system with large-scale weakly labeled data (e.g., AudioSet [31]), by first using a sound event detection model [32] to detect the anchor segment of sound events which are further used to constitute the artificial mixtures for the training of audio-queried sound separation models. These audio-queried sound separation approaches have shown great potential in the separation of unseen sound sources. However, during test time, the preparation of reference audio samples for the desired sound is often a time-consuming process.

3) Label-queried sound separation: An intuitive way to query a specific sound source is to use the label of its sound class [19]–[22]. Although acquiring a label query involves less effort than acquiring an audio or visual query, the label set is often pre-defined and is limited to a finite set of source categories. This imposes a challenge when attempting to generalize the separation system into an open-domain scenario, which may require re-training the sound separation model, or
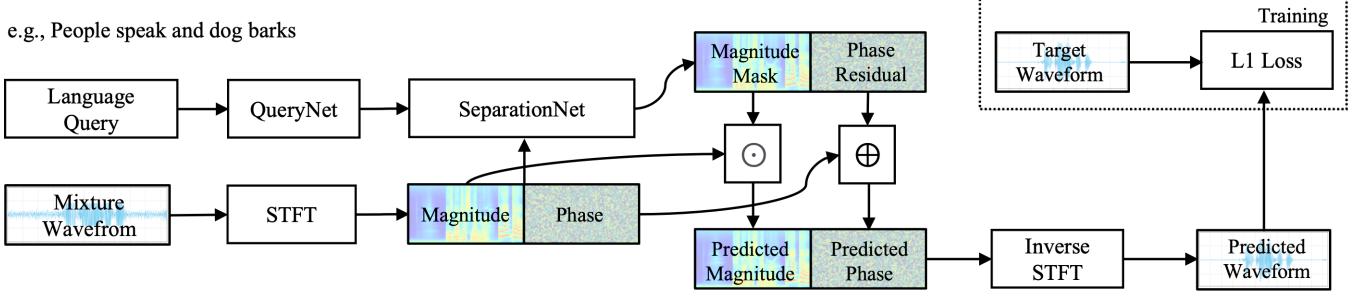
e.g., People speak and dog barks

Fig. 1. Framework of AudioSep. AudioSep has two key components: a QueryNet and a SeparationNet. The QueryNet is the text encoder of CLIP [34] or CLAP [35] model. The SeparationNet is a frequency-domain ResUNet [5], [15] model.

the use of complicated methods such as continual learning [20], [33]. Label information lacks the capability to describe the relationship between multiple sound events, such as their spatial relation and temporal order. This poses a challenge when the user intends to separate multiple sound sources rather than a single sound event.

III. AUDIOSEP

A. Language-queried audio source separation

Language-queried audio source separation (LASS) is a recently proposed new paradigm of QSS. LASS uses a natural language description of an arbitrary target source to separate the source from an audio mixture. Such natural language descriptions can include auxiliary information for describing the target source, such as spatial and temporal relationships of sound events, for example "*a dog barks in the background*" or "*people applaud followed by a woman speaking*".

LASS-Net [3] was our first attempt to perform end-to-end language-queried sound separation. LASS-Net consists of a language query encoder and a separation model. The separation model performs target source separation in the frequency domain and the target waveform is reconstructed using the noisy phase and inverse short-time Fourier transform (iSTFT). LASS-Net is trained on a subset (∼17.3 hours, 33 sound categories) of the AudioCaps [36] dataset and has shown great success in separating a wide range of sounds using audio caption queries. Kilgour et al. [24] proposed a similar model that accepts audio or text queries in a hybrid manner. Tzinis et al. [37] propose an optimal condition training (OCT) strategy for LASS. OCT performs greedy optimization toward the highest-performing condition among multiple conditions (e.g., signal energy, harmonicity) associated with a given target source, to improve the separation performance.

These LASS methods [3], [24], [37] require labeled text-audio paired data for supervised training, while such labeled data is often limited in practice. Recent work [23], [25] has investigated the potential of a self-supervised approach for LASS, including leveraging the visual modality as a bridge to learn the desired audio-textual correspondence. Instead of using a pre-trained language model (e.g., BERT [38]) [3], Dong et al. [23] use the contrastive language-image pretraining (CLIP) [34] model as the query encoder and train a LASS model conditioned on the visual context of

unlabeled noisy videos. Thanks to the aligned embedding space learned by the CLIP model, at the inference time, the separation model can be queried with text inputs in a zero-shot setting. Experimental results show that combining text and image conditions for hybrid training leads to better text-queried sound separation performance on musical instrument separation. Although preliminary studies have been undertaken into LASS, existing approaches work under the constraints of limited-source scenarios, such as musical instruments and a restricted set of universal sound classes. As such, these approaches have not met the expectation of LASS for zero-shot, open-domain sound separation.

B. Multimodal audio-language learning

Recently, the field of multi-modal audio-language has emerged as an important research area in audio signal processing and natural language processing. Audio-language tasks hold potential in various application scenarios. For instance, automatic audio captioning [39]–[46] aims to provide meaningful language descriptions of audio content, benefiting the hearing-impaired in comprehending environmental sounds. Language-based audio retrieval [47]–[51] facilitates efficient multimedia content retrieval and sound analysis for security surveillance. Text-to-audio generation [52]–[60] aims to synthesize audio content based on language descriptions, serving as sound synthesis tools for film-making, game design, virtual reality, and digital media, and aiding text understanding for the visually impaired. Contrastive language-audio pre-training (CLAP) [35] aims to learn an aligned audio-text embedding space via contrastive learning. CLAP facilitates downstream audio-text multimodal tasks (e.g., zero-shot audio classification) [52], [61]. In this work, we focus on the intersection between audio source separation and natural language processing, which is an important field for CASA research but less explored.

We next introduce AudioSep, a foundation model for open-domain sound separation with natural language queries. AudioSep has two key components: a QueryNet and a SeparationNet, as illustrated in Figure 1. We will introduce the details of each component below.

C. QueryNet

For QueryNet, we use the text encoder of the contrastive language-image pre-training model (CLIP) [34] or contrastive language-audio pre-training model (CLAP) [35]. The QueryNet is used to extract the text embedding of the natural language query. The input text query is denoted as $q = \{q_n\}_{n=1}^{N}$, consisting of a sequence of N word tokens. The sequence is processed by a text encoder to obtain the text embedding for the input language query. The text encoder encodes the input text tokens via a stack of Transformer blocks. After passing through the transformer layers, the output representations are aggregated, resulting in a fixed-length D-dimensional vector representation, where D corresponds to the latent dimension of the CLIP or CLAP. The text encoder is frozen during training.

CLIP is pre-trained on large-scale image-text paired data via contrastive learning. The CLIP text encoder learns to map textual descriptions to the same semantic space as the visual representations. The advantage of using CLIP text encoder in our task is, we are able to train or scale up the LASS model from large-scale unlabeled audio-visual data [23], by using visual embedding as an alternative to annotated audio-text paired data [23], [25]. Similar to CLIP, CLAP connects language and audio by using an audio encoder and a text encoder via a contrastive learning objective, bringing audio and text descriptions into a joint audio-text latent space. The CLAP text encoder may represent text that is better aligned with the audio context than the CLIP text encoder.

D. SeparationNet

For SeparationNet, we apply the frequency-domain ResUNet model [5], [15] as the separation backbone. The input to the ResUNet model is a mixture of audio clips. First, we apply a short-time Fourier transform (STFT) on the waveform to extract the complex spectrogram $X \in \mathbb{C}^{T \times F}$, the magnitude spectrogram and phase of X are denoted as $|X|$ and $e^{j\angle X}$, where $X = |X|e^{j\angle X}$. Then, we follow the same setting of [15], and we construct an encoder-decoder network to process the magnitude spectrogram. The ResUNet encoder-decoder comprises 6 encoder blocks, 4 bottleneck blocks, and 6 decoder blocks. In each encoder block, the spectrogram is downsampled into a bottleneck feature using 4 residual convolutional blocks, while each decoder block utilizes 4 residual deconvolutional blocks to upsample the feature and obtain the separation components. A skip connection is established between each encoder block and the corresponding decoder block, operating at the same downsampling/upsampling rate. The residual block consists of 2 CNN layers, 2 batch normalization layers, and 2 Leaky-ReLU activation layers. Furthermore, we introduce an additional residual shortcut connecting the input and output of each residual block. The ResUNet model inputs the complex spectrogram X and outputs the magnitude mask $|M|$ and the phase residual $\angle M$ conditioned on the text embedding e_q. $|M|$ controls how much the magnitude of $|X|$ should be scaled, and the angle $\angle M$ controls how much the angle of $\angle X$ should be rotated. The separated complex spectrogram can be obtained by multiplying the STFT of the mixture and the predicted magnitude mask $|M|$ and phase residual $\angle X$:

$$\hat{S} = |M| \odot |X|e^{j(\angle X + \angle M)}, \tag{1}$$

where $\odot$ is the Hadamard product.

To bridge the text encoder and the separation model, we use a Feature-wise Linearly modulated (FiLm) layer [62] after each ConvBlock deployed in the ResUNet. Specifically, let $H^{(l)} \in \mathbb{R}^{m \times h \times w}$ denote the output feature map produced by ConvBlock l with m channels, here h and w are the height and width of the feature map $H^{(l)}$, respectively. The modulation parameters are applied per feature map $H_i^{(l)}$ with the FiLm layer as follows:

$$\text{FiLM}(H_i^{(l)}|\gamma_i^{(l)}, \beta_i^{(l)}) = \gamma_i^{(l)} H_i^{(l)} + \beta_i^{(l)} \tag{2}$$

where $H_i^{(l)} \in \mathbb{R}^{h \times w}$, and $\gamma^{(l)}, \beta^{(l)} \in \mathbb{R}^m$ are the modulation parameters from $g(.)$, i.e., $(\gamma, \beta) = g(e_q)$, such that $g(.)$ is a neural network and e_q is the text embedding obtained from the text encoder. In this work, we model $g(\cdot)$ with two fully connected layers followed by ReLU activation, which is jointly trained with the ResUNet separation model.

E. Loss and training

During training, we use the loudness augmentation method proposed in [15]. When constituting the mixture x with s_1 and s_2, we first calculate the energy of s_1 and s_2 as E_1 and E_2 by $E = \|s\|_2^2$. We then apply a scaling factor $\alpha = \sqrt{E_1/E_2}$ to s_2 so that the s_1 and s_2 have the same energy:

$$x = s_1 + \alpha s_2. \tag{3}$$

We train AudioSep end-to-end using an L1 loss function between the predicted and target waveforms. Since waveform-based L1 loss is simple to implement and has shown good performance on universal sound separation tasks [15].

$$\text{Loss}_{L1} = \|s - \hat{s}\|_1 \tag{4}$$

The lower L1 loss value indicates that the separated signal $\hat{s}$ is closer to the ground truth signal s.

IV. DATASETS AND EVALUATION BENCHMARK

In this section, we provide a detailed description of the training dataset used for AudioSep, along with the established evaluation benchmark. The statistics of the dataset are presented in Tables I and II.

A. Training datasets

1) AudioSet: AudioSet [31] is a large-scale, weakly-labelled audio dataset with 2 million 10-second audio snippets sourced from YouTube. Each clip in the collection is categorised by the presented sound classes, without the timing information of sound events. AudioSet includes an ontology[1] of 527 distinct sound classes, such as "Human sounds", "Music", "Natural Sounds", among others. The training set comprises 2 063 839 clips, including a balanced subset of 22 160 audio

---

[1] https://research.google.com/audioset/ontology/index.html

clips. For each sound class in the balanced subset, there are at least 50 audio clips. After accounting for unavailable YouTube links, we were able to download $1\,934\,187$ audio clips with their corresponding video streams, equating to 94% of the complete training set. All the clips are either padded with silence or cut short to a duration of 10 seconds. Since a large amount of YouTube audio recordings have a sampling rate below 32 kHz, we have converted all audio recordings to a mono format and resampled them at 32 kHz. All audio clips within the training set are used as training data.

2) VGGSound: VGGSound [63] is a large-scale audio-visual dataset sourced from YouTube. VGGSound contains nearly $200\,000$ video clips each of length 10 seconds, annotated across 309 sound classes consisting of human actions, sound-emitting objects, and human-object interactions. The creation process of VGGSound has ensured that the object producing each sound is also discernible in the corresponding video clip. We utilize the original version of the VGGSound dataset, which includes $183\,727$ audio-visual clips for training and $15\,449$ for testing. We resample all audio clips at 32 kHz. All data within the training split are used for training.

3) AudioCaps: AudioCaps [36] is the largest publicly available audio captioning dataset, with $50\,725$ 10-second audio clips sourced from AudioSet. AudioCaps is divided into three splits: training, validation, and testing sets. The audio clips are annotated by humans with natural language descriptions through the Amazon Mechanical Turk crowd-sourced platform. Each audio clip in the training sets has a single human-annotated caption, while each clip in the validation and test set has five human-annotated captions. We retrieved AudioCaps based on the AudioSet we downloaded. The AudioCaps dataset we obtained consists of $49\,274$ out of $49\,837$ audio clips (98%) in the training set, 494 out of 495 clips (99%) in the validation set, and 957 out of 975 clips (98%) in the test set. All audio clips within the training and validation sets are used for our training process.

4) Clotho v2: Clotho v2 [64] is an audio captioning dataset that comprises sound clips obtained from the FreeSound platform[2]. Each audio clip in Clotho has been human-annotated via the Amazon Mechanical Turk crowd-sourced platform. Particular attention was paid to fostering diversity in the captions during the annotation process. In this work, we use Clotho v2 which was released for Task 6 of the DCASE 2021 Challenge[3]. Clotho v2 contains 3839, 1045, and 1045 audio clips for the development, validation, and evaluation split respectively. The sampling rate of all audio clips in the Clotho dataset is $44\,100$ Hz, each with five captions. Audio clips are of 15 to 30 seconds in duration and captions are 8 to 20 words long. We merge the development and validation split, forming a new training set with 4884 audio clips. All audio clips within the new training set and evaluation set are resampled at 32 kHz.

5) WavCaps: WavCaps [65] is a recently released large-scale weakly-labeled audio captioning dataset, comprising $403\,050$ audio clips with paired captions, totaling approxi-

mately 7568 hours. The audio clips constituting WavCaps originate from diverse sources, including FreeSound, BBC Sound Effects[4], SoundBible[5], and AudioSet. The audio captions are filtered and generated using the assistance of ChatGPT[6] based on the online-harvested raw audio descriptions. The average duration of audio clips is 67.59 seconds and the average text length of captions is 7.8 words. We resampled all audio clips within WavCaps at 32 kHz for training.

TABLE I
AUDIOSEP TRAINING DATASETS.

	Caption	Label	Video	Num. clips	Hours
AudioSet	×	✓	✓	$2\,063\,839$	5800
VGGSound	×	✓	✓	$183\,727$	550
AudioCaps	✓	✓	✓	$49\,768$	145
Clotho v2	✓	×	×	4884	37
WavCaps	✓	×	×	$40\,350$	7568

B. Evaluation benchmark

1) AudioSet: The evaluation set of AudioSet [31] contains $20\,317$ audio clips with 527 sound classes. We downloaded $18\,887$ audio clips from the evaluation set (93%) out of $20\,317$ audio clips. Source separation on AudioSet presents a considerable challenge, given that AudioSet contains a diverse range of sounds within a hierarchical ontology. To create evaluation data, we adopt the pipeline proposed in [15], which uses a sound event detection system [32] to analyze each 10-second audio clip to pinpoint anchor segments. Subsequently, two anchor segments from different sound classes are selected and combined to form a mixture with a signal-to-noise ratio (SNR) of 0 dB. We generate 10 mixtures for each sound class, leading to 5270 mixtures for all 527 sound classes in total.

2) VGGSound: In a similar way to [23], we manually selected 100 clean samples that each contain a distinct target sound event from the VGGSound test set. We refer to this set of 100 samples as VGGSound-Clean. For each audio sample from VGGSound-Clean, we randomly selected 10 audio samples from the remaining VGGSound test set to generate mixtures. Specifically, following [66], we first uniformly sampled the loudness of the two audio samples between -35 dB and -25 dB LUFS (Loudness Units Full Scale); then we mixed the signals together. The mixtures were scaled to 0.9 if clipping occurred. Finally, we constructed an evaluation set with 1000 samples. The average SNR of the evaluation set is around 0 dB.

3) AudioCaps: Our downloaded test set of the AudioCaps dataset [36] includes 957 audio clips, each annotated with five captions. To generate audio mixtures, we initially select an audio clip from the test set to serve as the target source, followed by a random selection of another audio clip as the background source, considering that the sound event tag[7] of the background source does not coincide with that of the target

---

source. For the test mixtures, each test audio is mixed with five randomly chosen background sources with an SNR at 0 dB. Each mixture is assigned one of the five audio captions of the target source. Consequently, 4785 test mixtures are created.

4) Clotho v2: The Clotho v2 [64] evaluation set includes 1045 audio clips, each provided with five human-annotated captions. The duration of audio clips varies between 15 and 30 seconds. For the creation of test mixtures, we designate each audio clip in the evaluation set as a target source. Subsequently, we select two audio clips at random from the evaluation set, concatenate them, and then truncate it to match the length of the target source, thereby producing the interference source. Applying this pipeline, each audio clip in the evaluation set is mixed with five audio clips at an SNR of 0 dB. Each created mixture is then assigned one of the five audio captions of the target source. This procedure culminates in a total of 5225 mixtures for evaluation.

5) ESC-50: The ESC-50 dataset [67] contains 2000 environmental audio recordings evenly arranged into 50 semantic classes including natural sounds, non-speech human sounds, domestic sounds, and urban noise. Each class contains 40 examples, with each audio clip having a duration of 5 seconds and with a sampling rate of 44.1 kHz. To make a consistent evaluation, we first downsample all audio clips at 32 kHz. Then, we randomly mix two audio clips from different sound classes with an SNR at 0 dB to form a pair. We constitute 40 mixtures for each sound class. This leads to a total of 2000 evaluation pairs, which are used to evaluate the zero-shot performance of our model on environmental sound separation.

6) MUSIC: The MUSIC dataset [12] is a collection of 536 video recordings of people playing a musical instrument from 11 instrument classes such as accordion, acoustic guitar, and cello. These video clips are crawled from YouTube and are relatively clean. Following the previous work [23], we downloaded the 46 video recordings in the test split. We further segmented all test videos into non-overlapping 10-second clips and resampled them at 32 kHz. For each video segment, we randomly select one segment from each of the other instrument classes to create a mixture with an SNR at 0 dB, resulting in a total of 5004 evaluation pairs, which are used to evaluate the zero-shot performance of our model on musical instrument separation.

7) Voicebank-DEMAND: The Voicebank-DEMAND dataset [68] integrates the Voicebank dataset [68], which includes clean speech, and the DEMAND [69] dataset, which encompasses a variety of background sounds that are used to create noisy speech. The noisy utterances are created by mixing the Voicebank dataset and the DEMAND dataset under signal-to-noise ratios of 15, 10, 5, and 0 dB. The test set of the Voicebank-DEMAND dataset includes a total of 824 utterances, which is used to evaluate the zero-shot performance of our model on speech enhancement. To make a fair comparison with previous speech enhancement systems [6], [7], [70], [71], we resample all audio clips to 16 kHz. We use "Speech" as the input text query to perform speech enhancement.

TABLE II
THE EVALUATION BENCHMARK FOR OPEN-DOMAIN SOUND SEPARATION WITH NATURAL LANGUAGE QUERIES.

	Num. mixtures	Sr (Hz)	Query Type
AudioSet	5270	32 k	text label
VGGSound	1000	32 k	text label
AudioCaps	4785	32 k	caption
Clotho v2	5225	32 k	caption
ESC-50	2000	32 k	text label
MUSIC	5004	32 k	text label
Voicebank-DEMAND	824	16 k	text label

C. Evaluation metrics

We utilize signal-to-distortion ratio improvement (SDRi) [15], [20] and scale-invariant SDR (SI-SDR) [72] to evaluate the performance of sound separation systems. For the speech enhancement task, following previous works [6], [7], [70], [71], we apply the perceptual evaluation of speech quality (PESQ) [73], mean opinion score (MOS) predictor of signal distortion (CSIG), MOS predictor of background-noise intrusiveness (CBAK), MOS predictor of overall signal quality (COVL) [74] and segmental signal-to-ratio noise (SSNR) [75] for evaluation. For each evaluation metric, higher values indicate better performance.

V. EXPERIMENTS

A. Training details

We randomly sample two audio segments from two audio clips from the training set and mix them together to constitute a training mixture. The length of the audio segment is 5 seconds. We extract the complex spectrogram from the waveform signal with a Hann window size of 1024 and a hop size of 320. We use the text encoder of CLIP [34] or CLAP [35] to extract the text embedding. For the CLIP model, we use the 'ViT-B-32' checkpoint. For the CLAP model, we use the publicly-available state-of-the-art checkpoint 'music_speech_audioset_epoch_15_esc_89.98.pt', which is trained on music, and speech datasets in addition to the original LAION-Audio-630k dataset [35]. The default configuration of AudioSep is using text supervision only. For processing video data, we uniformly extract frames with one-second intervals and compute their averaged CLIP embedding as the query embedding. For the separation model, we use a 30-layer ResUNet consisting of 6 encoder and 6 decoder blocks, which is the same as the previous work [15] on universal sound separation. Each encoder block consists of two convolutional layers with kernel sizes of 3×3. The number of output feature maps of the encoder blocks is 32, 64, 128, 256, 512, and 1024, respectively. The decoder blocks are symmetric to the encoder blocks. We apply an Adam optimizer with a learning rate of 1×10^{-3} to train the AudioSep with the batch size of 96. We train the AudioSep model for 1M steps on 8 Tesla V100 GPU cards.

B. Comparison systems

1) LASS models: We employ two state-of-the-art publicly available LASS models as the comparison systems. The first

TABLE III
BENCHAMRK EVALUATION RESULTS OF AUDIOSEP AND COMPARISON WITH BASELINE SYSTEMS.

	AudioSet		VGGSound		AudioCaps		Clotho		MUSIC		ESC-50		Voicebank-DEMAND	
	SI-SDR	SDRi	SI-SDR	SDRi	SI-SDR	SDRi	SI-SDR	SDRi	SI-SDR	SDRi	SI-SDR	SDRi	PESQ	SSNR
USS-ResUNet30 [15]	-	5.57	-	-	-	-	-	-	-	-	-	-	2.18	9.00
USS-ResUNet60 [15]	-	5.70	-	-	-	-	-	-	-	-	-	-	2.40	**9.35**
LASSNet [3]	-3.64	1.47	-4.50	1.17	-0.96	3.32	-3.42	2.24	-13.55	0.13	-2.11	3.69	1.39	0.98
CLIPSep [23]	-0.19	2.55	1.22	3.18	-0.09	2.95	-1.48	2.36	-0.37	2.50	-0.68	2.64	2.13	1.56
AudioSep-CLIP	**6.60**	**7.37**	7.24	7.50	5.95	7.45	4.54	6.28	**9.14**	**10.45**	8.90	10.03	2.40	8.09
AudioSep-CLAP	6.58	7.30	**7.38**	**7.55**	**6.45**	**7.68**	**4.84**	**6.51**	8.45	9.75	**9.16**	**10.24**	**2.41**	8.95

one is LASS-Net [3], which uses a pre-trained BERT and ResUNet as the text query encoder and the separation model, respectively. LASS-Net is trained on a subset ($\sim$17.3 hours) of AudioCaps including universal sounds of categories such as human sounds, animal, sounds of things, natural sounds, and environmental sounds. The second one is CLIPSep, which uses CLIP [34] as the query encoder and a model based on Sound-of-Pixels (SOP) [12] for separation. CLIPSep [23] is trained with noisy audio-visual videos ($\sim$500 hours) from VGGSound [63] dataset using hybrid vision and text supervision signal. In addition, CLIPSep employs a training strategy called noise invariant training (NIT), designed to enable robust learning from noisy video data. As the CLIPSep model is trained with 16 kHz sampling rate, we downsample the audio mixtures and corresponding ground truth signal to 16 kHz for evaluation. Both LASS-Net and CLIPSep are frequency domain separation models, the noisy phase is used to reconstruct the waveform in the test time.

2) Audio-queried sound separation models: We use audio-queried separation models as comparison systems. Kong et al. [15] proposed a universal sound separation system trained on AudioSet. This system performs query-based sound separation by using the average audio embedding calculated from query examples from the training set as the condition and can separate hundreds of sound classes using a single model. Empirical studies [15] conducted by Kong et al. have assessed the effectiveness of various sound separation systems such as ConvTasNet [9], UNet [76], and ResUNet [5], along with different audio embedding extractors such as PANNs [32] and HTSAT [77]. The results indicate that a system combining PANNs and ResUNet achieved the best performance. As a result, we adopt the PANN-ResUNet separation model with 30 and 60 layers of the ResUNet as the comparison systems, denoted as USS-ResUNet30 and USS-ResUNet60, respectively.

3) Speech enhancement models: Following [7], we adopt four off-the-shelf speech enhancement models as the comparison systems including Wiener filter [70], SEGAN [6], AudioSet-UNet [7], and Wave-U-Net [71]. Wiener filter [70] is a method based on signal processing methods. SEGAN is designed with the generative adversarial network [78]. AudioSet-UNet is a frequency-domain UNet-based model trained with weakly labeled AudioSet [31] data. Wave-U-Net is a time-domain UNet-based speech enhancement model.

TABLE IV
SPEECH ENHANCEMENT RESULTS ON VOICEBANK-DEMAND DATASET.

	PESQ	CSIG	CBAK	COVL	SSNR
Noisy	1.97	3.35	2.44	2.63	1.68
Wiener [70]	2.22	3.23	2.68	2.67	5.07
SEGAN [6]	2.16	3.48	2.94	2.80	7.73
AudioSet-UNet [7]	2.28	2.43	2.96	2.30	8.75
Wave-U-Net [71]	2.40	**3.52**	**3.24**	**2.96**	**9.97**
AudioSep-CLIP	2.40	3.37	3.09	2.88	8.09
AudioSep-CLAP	**2.41**	3.30	3.15	2.84	8.95

C. Evaluation results on seen datasets

We first assess the performance of AudioSep on datasets seen during training including AudioSet, VGGSound, AudioCaps, and Clotho, as shown in Table III. The AudioSep model with CLIP and CLAP text encoders are referred to as AudioSep-CLIP and AudioSep-CLAP, respectively. AudioSep shows strong sound separation performance using text labels or audio captions as input queries. On the AudioSet, AudioSep-CLIP achieves an SI-SDR and SDRi of 6.6 dB and 7.37 dB across 527 audio event classes, respectively. In comparison, AudioSep-CLAP achieved an SI-SDR of 6.58 dB and an SDRi of 7.3 dB. For the VGGSound dataset, the AudioSep-CLAP model demonstrated an SI-SDR of 7.38 dB and an SDRi of 7.55 dB. On the AudioCaps and Clotho datasets, the AudioSep-CLAP model achieved an SI-SDR of 6.45 dB and 4.84 dB, respectively, along with SDRi values of 7.68 dB and 6.51 dB, respectively. Overall, the AudioSep-CLIP model under-performs AudioSep-CLAP on VGGSound, AudioCaps, and Clotho datasets, suggesting that the CLAP text embedding outperforms CLIP text embedding for universal sound separation in our experiments.

Two state-of-the-art models LASS-Net and CLIPSep did not perform well on the separation of target sounds on our benchmark, since they are trained using audio data with limited concepts. Audio-queried sound separation baseline systems USS-ResUNet30 and USS-ResUNet60 have achieved an SDRi of 5.57 dB and 5.7 dB, respectively, which underperform AudioSep. Notably, in the inference stage, AudioSep only requires the text description of the target source, which is more convenient to obtain, as compared with anchor audio required in the audio-queried methods. Evaluation results on seen datasets indicate the superior performance of AudioSep compared with previous state-of-the-art LASS models and off-the-shelf audio-queried sound separation models.

TABLE V
ABLATION STUDY OF SCALING UP AUDIOSEP WITH MULTIMODAL SUPERVISION.

	AudioSet		VGGSound		AudioCaps		Clotho		MUSIC		ESC-50		Voicebank-DEMAND	
	SI-SDR	SDRi	SI-SDR	SDRi	SI-SDR	SDRi	SI-SDR	SDRi	SI-SDR	SDRi	SI-SDR	SDRi	PESQ	SSNR
AudioSep-CLIP-TR0.0	0.45	2.91	-3.84	0.78	-0.02	3.29	-1.50	2.42	-0.21	2.97	-0.24	3.67	2.27	6.67
AudioSep-CLIP-TR0.5	6.48	7.28	7.01	7.27	5.84	7.25	4.32	6.10	7.40	9.05	8.74	9.97	2.37	**9.07**
AudioSep-CLIP-TR0.75	6.52	7.31	**7.46**	**7.67**	**5.98**	7.41	4.38	6.16	8.00	9.35	**8.92**	**10.10**	**2.40**	8.83
AudioSep-CLIP-TR1.0	**6.60**	**7.37**	7.24	7.50	5.95	**7.45**	**4.54**	**6.28**	**9.14**	**10.45**	8.90	10.03	**2.40**	8.09
AudioSep-CLAP-TR0.0	0.95	3.35	-3.01	0.88	-1.10	2.57	-1.45	2.31	-1.36	2.36	-1.95	2.74	2.32	6.42
AudioSep-CLAP-TR0.5	5.52	6.46	5.40	5.93	5.43	6.89	3.98	5.84	6.80	8.47	6.77	8.42	2.40	8.82
AudioSep-CLAP-TR0.75	6.13	6.96	6.88	7.07	6.02	7.35	4.55	6.27	7.80	9.24	8.48	9.69	2.45	8.30
AudioSep-CLAP-TR1.0	**6.58**	**7.30**	**7.38**	**7.55**	**6.45**	**7.68**	**4.84**	**6.51**	**8.45**	**9.75**	**9.16**	**10.24**	**2.41**	**8.95**

D. Evaluation results on unseen datasets

We conducted further evaluations to assess the zero-shot separation performance of AudioSep on unseen datasets, including MUSIC, ESC-50, and Voicebank-DEMAND. AudioSep achieved promising zero-shot separation performance on MUSIC musical instrument separation and ESC-50 sound effects separation, as shown in Table IV. Particularly, on MUSIC, AudioSep-CLIP achieved an SI-SDR of 9.14 dB and an SDRi of 10.45 dB, considerably outperforming AudioSep-CLAP which achieved an SI-SDR of 8.45 dB and an SDRi of 9.75 dB. The experimental results indicate that the CLIP text embedding captures the difference between musical instruments better than the CLAP text embedding. For ESC-50, AudioSep CLAP achieved an SI-SDR of 9.16 dB and an SDRi of 10.24 dB, which slightly outperforms the AudioSep-CLIP which achieved an SI-SDR of 8.9 dB and an SDRi of 10.03 dB. Although CLIPSep achieved good performance in separating musical instruments from background noise [23], its performance degrades when faced with musical instrument mixtures. Neither CLIPSep nor LASS-Net performs well in these evaluation datasets.

Table V shows the results of Voicebank-DEMAND speech enhancement. Noisy speech without enhancement has PESQ, CSIG, CBAK, COVL, and SSNR of 1.97 dB, 3.35 dB, 2.44 dB, 2.63 dB and 1.68 dB respectively. AudioSep-CLIP results of PESQ, CSIG, CBAK, COVL, and SSNR values of 2.40 dB, 3.37 dB, 3.09 dB, 2.88 dB and 8.09 dB respectively. AudioSep-CLAP achieves PESQ, CSIG, CBAK, COVL, and SSNR of 2.41 dB, 3.30 dB, 3.15 dB, 2.84 dB and 8.95 dB respectively. AudioSep-CLAP surpasses all the speech enhancement baselines in the PESQ metric. In addition, both AudioSep-CLIP and AudioSep-CLAP outperform the Wiener [70] and AudioSet-UNet [7] method across all metrics, showing the good speech enhancement capability of the AudioSep model. On the other hand, AudioSep performs less effectively than the SEGAN [6] method in the CSIG metric, indicating that AudioSep may lose certain speech details, particularly in the high-frequency component. This discrepancy could potentially be due to the presence of large-scale narrow-band signals in the AudioSet training data.

Experimental results on the unseen datasets indicate the promising zero-shot generalization performance of AudioSep. With AudioSep, we may perform sound separation tasks on new data distributions with textual queries.

E. Leveraging multimodal supervision for AudioSep

Recent research has explored the potential of utilizing multimodal supervision [23]–[25] to enhance the scalability of training LASS models. However, these works mainly focused on small-scale training sets. In this section, we present ablation studies to investigate the efficacy of using large-scale multimodal supervision with CLIP and CLAP models to scale up AudioSep. We aim to gain insights into the applicability and performance of leveraging large-scale multimodal supervision for LASS.

For the CLIP model, we use hybrid vision-language supervision for AudioSep with varying text ratios (TR). Specifically, we experiment with text ratios of 0%, 50%, 75%, and 100%. The resulting AudioSep models are referred to as AudioSep-CLIP-TR0.0, AudioSep-CLIP-TR0.5, AudioSep-CLIP-TR0.75, and AudioSep-CLIP-TR1.0, respectively. The hybrid supervision approach is exclusively adopted for the training process using AudioSet [31] and VGGSound [63] datasets that have video stream information. For datasets that solely consist of audio and text, such as WavCaps [65], we continue to use text supervision. This training configuration allows us to leverage the available modalities and tailor the supervision approach to suit the specific characteristics of each dataset. Experimental results are shown in the upper part of Table V. When training AudioSep-CLIP0.0 without text supervision using the AudioSet and VGGsound datasets, the performance is clearly inadequate across all evaluation datasets. As large-scale video data may contain irrelevant audio contexts, this experimental result highlights the key of text supervision from audio event labels provided by AudioSet and VGGSound in effectively training AudioSep. When training AudioSep using text ratios of 50% and 75%, the overall performance is comparable to AudioSep-CLIP-TR1.0 trained exclusively with text supervision. This observation suggests that incorporating additional supervision from the visual modality does not enhance the separation performance in large-scale training settings, which can potentially be attributed to the inherent noise within video data at scale.

For the CLAP setting, similar to CLIP, we employ a hybrid audio-text supervision approach for AudioSep. We conduct experiments with text ratios of 0%, 50%, 75%, and 100%. The resulting AudioSep models are denoted as AudioSep-CLAP-TR0.0, AudioSep-CLAP-TR0.5, AudioSep-CLAP-TR0.75, and AudioSep-CLAP-TR1.0. Experimental results are shown in the bottom part of Table V. In contrast to

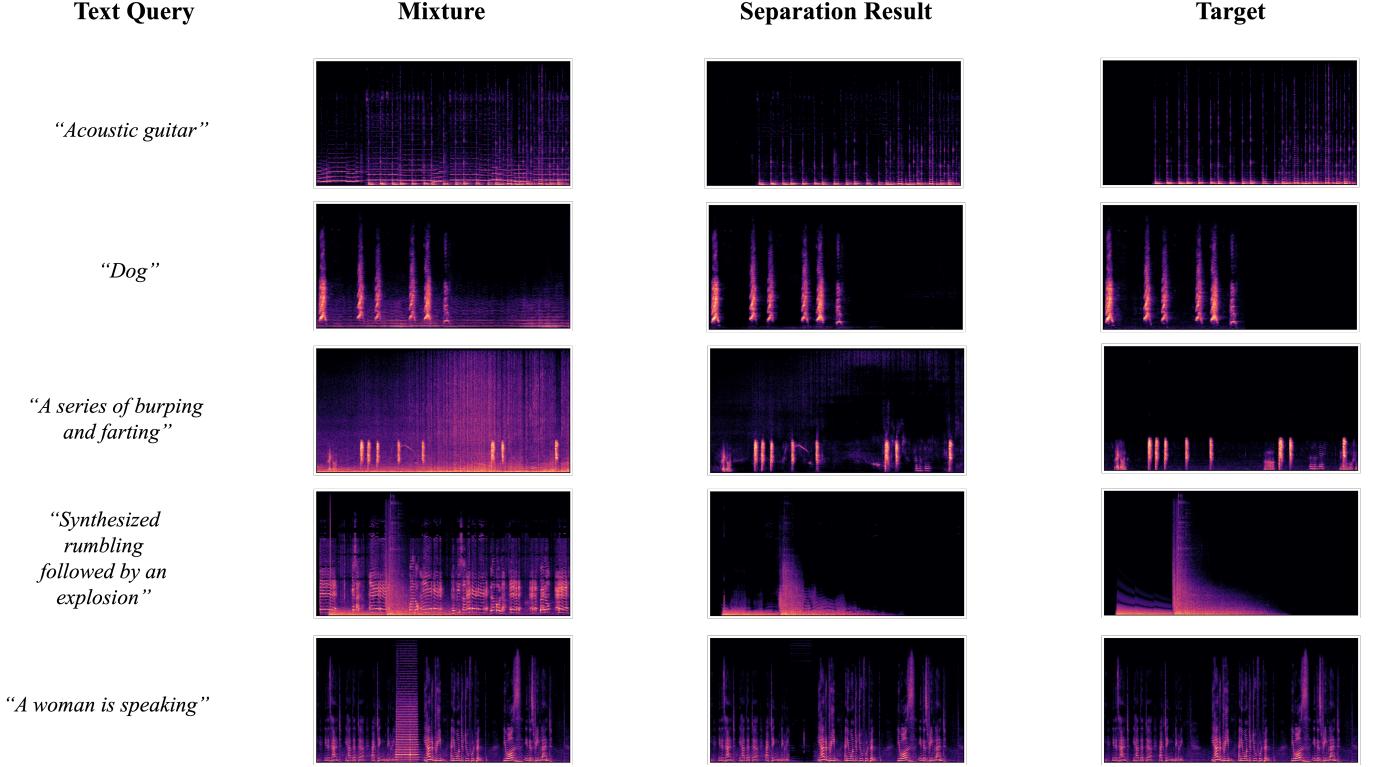| Text Query	Mixture	Separation Result	Target

Fig. 2. Visualization of separation results obtained by the AudioSep-CLAP model.

the findings in the CLIP setting, utilizing hybrid audio-text supervision for training results in a decrease in separation performance. Specifically, we observed that the inclusion of more audio supervision leads to lower separation performance. Given the substantial difference in training data size, with CLIP trained on a dataset of around 400M image-text pairs and CLAP on approximately 630K audio-text pairs, we suggest the reason could be due to the audio-text latent spaces of CLAP not being aligned well. Consequently, AudioSep may tend to overfit the audio latent when a large amount of audio supervision is utilized, which further results in performance degradation using text queries in the inference stage.

F. Visualization of separation results

We performed visualizations of spectrograms for audio mixtures, and target audio sources, and separated sources using text queries of diverse sound sources (e.g., musical instrument, audio event, speech) with the AudioSep-CLAP model, as shown in Figure 2. We observe that the spectrogram pattern of the separated source is close to the target source, which is consistent with our objective experimental results. Audio samples are available on our project page[8].

G. Comparison of various text queries

In practice, human descriptions of an audio source are generally personalized. For this purpose, we randomly selected 50 audio clips from the test set of the AudioCaps dataset.

[8]https://audio-agi.github.io/Separate-Anything-You-Describe

We engaged four language experts from the University of Surrey to individually annotate the selected clips. Each expert provided a single description per audio clip without any specific hints or restrictions. Consequently, we obtained a set of six descriptions for each audio clip. This set included one caption sourced from AudioCaps, four captions provided by the language experts, and the AudioSet audio event text labels. To create test mixtures, each audio clip was mixed with ten background sources randomly selected with an SNR at 0 dB, considering that the sound event labels of the background source do not overlap with that of the target source. We denote this test dataset as AudioCaps-Mini.

We evaluate the performance of AudioSep-CLIP and AudioSep-CLAP on AudioCaps-Mini. To investigate the effects of various text queries, we utilize the retrieved AudioSet event labels, the original AudioCaps audio captions, and our re-annotated natural language descriptions, which are referred to as the "text label" "original caption" and "re-annotated caption", respectively. An example of AudioCaps-Mini can be found in Table VI. Experimental results are shown in Table VII. For both the AudioSep-CLIP and AudioSep-CLAP models, we have observed a considerable performance improvement when utilizing the "original caption" as text queries instead of using the "text label". This may be attributed to the fact that human-annotated captions provide more comprehensive and accurate descriptions of the source of interest compared to audio event labels. Despite the personalized nature and different word distribution of our re-annotated captions, the results obtained using the "re-annotated caption" are slightly worse than those using the "original caption",

while still marginally outperforming the results obtained with the "text label". These experimental findings demonstrate the promising generalization performance and robustness of AudioSep in real-world cases with diverse language queries.

TABLE VI
AN EXAMPLE FROM AUDIOCAPS-MINI.

	Text Description
Text label	"Vibration"
Original caption	"Clicking followed by vibrations"
Re-annotated caption1	"Gear change, moving vehicle"
Re-annotated caption2	"The distant engine sound"
Re-annotated caption3	"The engine is running in distant"
Re-annotated caption4	"The engine is starting"

TABLE VII
EXPERIMENTAL RESULTS OF COMPARISON OF DIFFERENT TEXT QUERIES
ON AUDIOCAPS-MINI DATASET.

	SI-SDR	SDRi
AudioSep-CLIP (text label)	6.39	7.70
AudioSep-CLIP (original caption)	7.27	8.25
AudioSep-CLIP (re-annotated caption)	6.44	7.75
AudioSep-CLAP (text label)	6.32	7.65
AudioSep-CLAP (original caption)	7.73	8.47
AudioSep-CLAP (re-annotated caption)	6.59	7.80

VI. CONCLUSION AND FUTURE WORK

We have introduced AudioSep, a foundation model for open-domain universal sound separation with natural language descriptions. AudioSep can perform zero-shot separation using text labels or audio captions as queries. We have presented a comprehensive evaluation benchmark including numerous sound separation tasks such as audio event separation, musical instrument separation, and speech enhancement. AudioSep outperforms state-of-the-art text-queried separation systems and off-the-shelf audio-queried sound separation models. We show that AudioSep is a promising approach to flexibly address the CASA problem with strong sound separation performance. In future work, we will improve the separation performance of AudioSep via unsupervised learning techniques [14], [27] and extend AudioSep to support vision-queried separation, audio-queried separation, and text-guided speaker separation [79] tasks.

ACKNOWLEDGMENTS

REFERENCES

[1] T. Virtanen, M. D. Plumbley, and D. Ellis, *Computational Analysis of Sound Scenes and Events*. Springer, 2018.

[2] S. Haykin and Z. Chen, "The cocktail party problem," *Neural Computation*, vol. 17, no. 9, pp. 1875–1902, 2005.

[3] X. Liu, H. Liu, Q. Kong, X. Mei, J. Zhao, Q. Huang, M. D. Plumbley, and W. Wang, "Separate what you describe: Language-queried audio source separation," in *INTERSPEECH*, 2022.

[4] I. Kavalerov, S. Wisdom, H. Erdogan, B. Patton, K. Wilson, J. Le Roux, and J. R. Hershey, "Universal sound separation," in *2019 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, 2019, pp. 175–179.

[5] Q. Kong, Y. Cao, H. Liu, K. Choi, and Y. Wang, "Decoupling magnitude and phase estimation with deep ResUNet for music source separation." in *ISMIR*, 2021.

[6] S. Pascual, A. Bonafonte, and J. Serra, "SEGAN: Speech enhancement generative adversarial network," in *INTERSPEECH*, 2017.

[7] Q. Kong, H. Liu, X. Du, L. Chen, R. Xia, and Y. Wang, "Speech enhancement with weakly labelled data from AudioSet," in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*, 2021.

[8] D. Wang and J. Chen, "Supervised speech separation based on deep learning: An overview," *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 26, no. 10, pp. 1702–1726, 2018.

[9] Y. Luo and N. Mesgarani, "Conv-TasNet: Surpassing ideal time–frequency magnitude masking for speech separation," *IEEE/ACM transactions on Audio, Speech, and Language Processing*, vol. 27, no. 8, pp. 1256–1266, 2019.

[10] S. Rubin, F. Berthouzoz, G. J. Mysore, W. Li, and M. Agrawala, "Content-based tools for editing audio stories," in *Proceedings of the 26th Annual ACM Symposium on User Interface Software and Technology*, 2013, pp. 113–122.

[11] Y. Peng, X. Huang, and Y. Zhao, "An overview of cross-media retrieval: Concepts, methodologies, benchmarks, and challenges," *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 28, no. 9, pp. 2372–2385, 2017.

[12] H. Zhao, C. Gan, A. Rouditchenko, C. Vondrick, J. McDermott, and A. Torralba, "The sound of pixels," in *Proceedings of the European Conference on Computer Vision (ECCV)*, 2018, pp. 570–586.

[13] M. Chatterjee, J. Le Roux, N. Ahuja, and A. Cherian, "Visual scene graphs for audio source separation," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021, pp. 1204–1213.

[14] E. Tzinis, S. Wisdom, A. Jansen, S. Hershey, T. Remez, D. P. Ellis, and J. R. Hershey, "Into the wild with AudioScope: Unsupervised audio-visual separation of on-screen sounds," *arXiv:2011.01143*, 2020.

[15] Q. Kong, K. Chen, H. Liu, X. Du, T. Berg-Kirkpatrick, S. Dubnov, and M. D. Plumbley, "Universal source separation with weakly labelled data," *arXiv:2305.07447*, 2023.

[16] K. Chen, X. Du, B. Zhu, Z. Ma, T. Berg-Kirkpatrick, and S. Dubnov, "Zero-shot audio source separation through query-based learning from weakly-labeled data," in *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, vol. 36, no. 4, 2022, pp. 4441–4449.

[17] B. Gfeller, D. Roblek, and M. Tagliasacchi, "One-shot conditional audio filtering of arbitrary sounds," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2021, pp. 501–505.

[18] Y. Wang, D. Stoller, R. M. Bittner, and J. P. Bello, "Few-shot musical source separation," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2022, pp. 121–125.

[19] B. Veluri, J. Chan, M. Itani, T. Chen, T. Yoshioka, and S. Gollakota, "Real-time target sound extraction," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023, pp. 1–5.

[20] M. Delcroix, J. B. Vázquez, T. Ochiai, K. Kinoshita, Y. Ohishi, and S. Araki, "Soundbeam: Target sound extraction conditioned on sound-class labels and enrollment clues for increased performance and continuous learning," *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 31, pp. 121–136, 2022.

[21] T. Ochiai, M. Delcroix, Y. Koizumi, H. Ito, K. Kinoshita, and S. Araki, "Listen to what you want: Neural network-based universal sound selector," *arXiv:2006.05712*, 2020.

[22] H. Wang, D. Yang, C. Weng, J. Yu, and Y. Zou, "Improving target sound extraction with timestamp information," *arXiv:2204.00821*, 2022.

[23] H.-W. Dong, N. Takahashi, Y. Mitsufuji, J. McAuley, and T. Berg-Kirkpatrick, "CLIPSep: Learning text-queried sound separation with noisy unlabeled videos," in *Proceedings of International Conference on Learning Representations (ICLR)*, 2023.

[24] K. Kilgour, B. Gfeller, Q. Huang, A. Jansen, S. Wisdom, and M. Tagliasacchi, "Text-driven separation of arbitrary sounds," in *INTERSPEEH*, 2022.

[25] R. Tan, A. Ray, A. Burns, B. A. Plummer, J. Salamon, O. Nieto, B. Russell, and K. Saenko, "Language-guided audio-visual source separation via trimodal consistency," *arXiv:2303.16342*, 2023.

[26] D. Yu, M. Kolbæk, Z.-H. Tan, and J. Jensen, "Permutation invariant training of deep models for speaker-independent multi-talker speech separation," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2017, pp. 241–245.

[27] S. Wisdom, E. Tzinis, H. Erdogan, R. Weiss, K. Wilson, and J. Hershey, "Unsupervised sound separation using mixture invariant training," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 3846–3857, 2020.

[28] R. Gao and K. Grauman, "VisualVoice: Audio-visual speech separation with cross-modal consistency," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 15 490–15 500.

[29] J. Wu, Y. Xu, S.-X. Zhang, L.-W. Chen, M. Yu, L. Xie, and D. Yu, "Time domain audio visual speech separation," in *IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*, 2019, pp. 667–673.

[30] J. H. Lee, H.-S. Choi, and K. Lee, "Audio query-based music source separation," *arXiv preprint arXiv:1908.06593*, 2019.

[31] J. F. Gemmeke, D. P. Ellis, D. Freedman, A. Jansen, W. Lawrence, R. C. Moore, M. Plakal, and M. Ritter, "Audio set: An ontology and human-labeled dataset for audio events," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2017, pp. 776–780.

[32] Q. Kong, Y. Cao, T. Iqbal, Y. Wang, W. Wang, and M. D. Plumbley, "PANNs: Large-scale pretrained audio neural networks for audio pattern recognition," *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 28, pp. 2880–2894, 2020.

[33] Y. Xiao, X. Liu, J. King, A. Singh, E. S. Chng, M. D. Plumbley, and W. Wang, "Continual learning for on-device environmental sound classification," in *Proceedings of the 7th Detection and Classification of Acoustic Scenes and Events (DCASE) Workshop*, 2022.

[34] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark *et al.*, "Learning transferable visual models from natural language supervision," in *International Conference on Machine Learning*, 2021, pp. 8748–8763.

[35] Y. Wu, K. Chen, T. Zhang, Y. Hui, T. Berg-Kirkpatrick, and S. Dubnov, "Large-scale contrastive language-audio pretraining with feature fusion and keyword-to-caption augmentation," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023, pp. 1–5.

[36] C. D. Kim, B. Kim, H. Lee, and G. Kim, "Audiocaps: Generating captions for audios in the wild," in *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, 2019, pp. 119–132.

[37] E. Tzinis, G. Wichern, P. Smaragdis, and J. Le Roux, "Optimal condition training for target source separation," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023, pp. 1–5.

[38] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of deep bidirectional transformers for language understanding," *arXiv:1810.04805*, 2018.

[39] X. Mei, X. Liu, Q. Huang, M. D. Plumbley, and W. Wang, "Audio captioning transformer," in *Proceedings of the 6th Detection and Classification of Acoustic Scenes and Events (DCASE) Workshop*, 2021, pp. 211–215.

[40] K. Drossos, S. Adavanne, and T. Virtanen, "Automated audio captioning with recurrent neural networks," in *IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, 2017, pp. 374–378.

[41] X. Mei, X. Liu, M. D. Plumbley, and W. Wang, "Automated audio captioning: An overview of recent progress and new challenges," *EURASIP Journal on Audio, Speech, and Music Processing*, vol. 2022, no. 1, pp. 1–18, 2022.

[42] X. Liu, Q. Huang, X. Mei, T. Ko, H. Tang, M. D. Plumbley, and W. Wang, "CL4AC: A contrastive loss for audio captioning," in *Proceedings of the 6th Detection and Classification of Acoustic Scenes and Events (DCASE) Workshop*, 2021, pp. 196–200.

[43] X. Liu, X. Mei, Q. Huang, J. Sun, J. Zhao, H. Liu, M. D. Plumbley, V. Kilic, and W. Wang, "Leveraging pre-trained BERT for audio captioning," in *30th European Signal Processing Conference (EUSIPCO)*, 2022, pp. 1145–1149.

[44] M. Kim, K. Sung-Bin, and T.-H. Oh, "Prefix tuning for automated audio captioning," in *ICASSP 2023-2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2023, pp. 1–5.

[45] Y. Koizumi, R. Masumura, K. Nishida, M. Yasuda, and S. Saito, "A transformer-based audio captioning model with keyword estimation," *arXiv:2007.00222*, 2020.

[46] X. Liu, Q. Huang, X. Mei, H. Liu, Q. Kong, J. Sun, S. Li, T. Ko, Y. Zhang, L. H. Tang *et al.*, "Visually-aware audio captioning with adaptive audio-visual attention," in *INTERSPEEH*, 2023.

[47] X. Mei, X. Liu, J. Sun, M. Plumbley, and W. Wang, "On metric learning for audio-text cross-modal retrieval," in *Proc. Interspeech*, 2022, pp. 4142–4146.

[48] X. Mei, X. Liu, H. Liu, J. Sun, M. D. Plumbley, and W. Wang, "Language-based audio retrieval with pre-trained models," *Detection and Classification of Acoustic Scenes and Events (DCASE) Challenge, Tech. Rep*, 2022.

[49] A.-M. Oncescu, A. Koepke, J. F. Henriques, Z. Akata, and S. Albanie, "Audio retrieval with natural language queries," *arXiv:2105.02192*, 2021.

[50] A. S. Koepke, A.-M. Oncescu, J. Henriques, Z. Akata, and S. Albanie, "Audio retrieval with natural language queries: A benchmark study," *IEEE Transactions on Multimedia*, 2022.

[51] S. Lou, X. Xu, M. Wu, and K. Yu, "Audio-text retrieval in context," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2022, pp. 4793–4797.

[52] H. Liu, Z. Chen, Y. Yuan, X. Mei, X. Liu, D. Mandic, W. Wang, and M. D. Plumbley, "AudioLDM: Text-to-audio generation with latent diffusion models," *arXiv:2301.12503*, 2023.

[53] Y. Yuan, H. Liu, J. Liang, X. Liu, M. D. Plumbley, and W. Wang, "Leveraging pre-trained audioldm for sound generation: A benchmark study," *arXiv:2303.03857*, 2023.

[54] D. Ghosal, N. Majumder, A. Mehrish, and S. Poria, "Text-to-audio generation using instruction-tuned llm and latent diffusion model," *arXiv:2304.13731*, 2023.

[55] J. Huang, Y. Ren, R. Huang, D. Yang, Z. Ye, C. Zhang, J. Liu, X. Yin, Z. Ma, and Z. Zhao, "Make-An-Audio 2: Temporal-enhanced text-to-audio generation," *arXiv:2305.18474*, 2023.

[56] D. Yang, J. Yu, H. Wang, W. Wang, C. Weng, Y. Zou, and D. Yu, "Diffsound: Discrete diffusion model for text-to-sound generation," *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 2023.

[57] F. Kreuk, G. Synnaeve, A. Polyak, U. Singer, A. Défossez, J. Copet, D. Parikh, Y. Taigman, and Y. Adi, "AudioGen: Textually guided audio generation," *arXiv preprint arXiv:2209.15352*, 2022.

[58] X. Liu, T. Iqbal, J. Zhao, Q. Huang, M. D. Plumbley, and W. Wang, "Conditional sound generation using neural discrete time-frequency representation learning," in *2021 IEEE 31st International Workshop on Machine Learning for Signal Processing (MLSP)*, 2021, pp. 1–6.

[59] Y. Yuan, H. Liu, X. Liu, X. Kang, P. Wu, M. D. Plumbley, and W. Wang, "Text-driven Foley sound generation with latent diffusion model," *arXiv:2306.10359*, 2023.

[60] X. Liu, Z. Zhu, H. Liu, Y. Yuan, M. Cui, Q. Huang, J. Liang, Y. Cao, Q. Kong, M. D. Plumbley *et al.*, "WavJourney: Compositional audio creation with large language models," *arXiv preprint arXiv:2307.14335*, 2023.

[61] J. Liang, X. Liu, H. Liu, H. Phan, E. Benetos, M. D. Plumbley, and W. Wang, "Adapting language-audio models as few-shot audio learners," in *INTERSPEEH*, 2023.

[62] E. Perez, F. Strub, H. De Vries, V. Dumoulin, and A. Courville, "FiLm: Visual reasoning with a general conditioning layer," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, no. 1, 2018.

[63] H. Chen, W. Xie, A. Vedaldi, and A. Zisserman, "VGGSound: A large-scale audio-visual dataset," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2020, pp. 721–725.

[64] K. Drossos, S. Lipping, and T. Virtanen, "Clotho: An audio captioning dataset," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2020, pp. 736–740.

[65] X. Mei, C. Meng, H. Liu, Q. Kong, T. Ko, C. Zhao, M. D. Plumbley, Y. Zou, and W. Wang, "WavCaps: A ChatGPT-assisted weakly-labelled audio captioning dataset for audio-language multimodal research," *arXiv:2303.17395*, 2023.

[66] J. Cosentino, M. Pariente, S. Cornell, A. Deleforge, and E. Vincent, "Librimix: An open-source dataset for generalizable speech separation," *arXiv:2005.11262*, 2020.

[67] K. J. Piczak, "ESC: Dataset for environmental sound classification," in *Proceedings of the 23rd ACM International Conference on Multimedia*, 2015, pp. 1015–1018.

[68] C. Veaux, J. Yamagishi, and S. King, "The voice bank corpus: Design, collection and data analysis of a large regional accent speech database," in *International Conference Oriental COCOSDA with Conference on Asian Spoken Language Research and Evaluation (O-COCOSDA/CASLRE)*, 2013, pp. 1–4.

[69] J. Thiemann, N. Ito, and E. Vincent, "DEMAND: a collection of multi-channel recordings of acoustic noise in diverse environments," in *Proceedings of Meetings on Acoustics*, 2013, pp. 1–6.

[70] P. Scalart *et al.*, "Speech enhancement based on a priori signal to noise estimation," in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*, vol. 2, 1996, pp. 629–632.

[71] C. Macartney and T. Weyde, "Improved speech enhancement with the Wave-U-Net," *arXiv:1811.11307*, 2018.

[72] J. Le Roux, S. Wisdom, H. Erdogan, and J. R. Hershey, "SDR–half-baked or well done?" in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2019, pp. 626–630.

[73] A. W. Rix, J. G. Beerends, M. P. Hollier, and A. P. Hekstra, "Perceptual evaluation of speech quality (PESQ) - a new method for speech quality assessment of telephone networks and codecs," in *IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP).*, vol. 2, 2001, pp. 749–752.

[74] Y. Hu and P. C. Loizou, "Evaluation of objective quality measures for speech enhancement," *IEEE Transactions on Audio, Speech, and Language Processing*, vol. 16, no. 1, pp. 229–238, 2007.

[75] S. R. Quackenbush, T. P. Barnwell, and M. A. Clements, "Objective measures of speech quality," *Prentice Hall*, 1988.

[76] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional networks for biomedical image segmentation," in *Medical Image Computing and Computer-Assisted Intervention–MICCAI 2015: 18th International Conference, Munich, Germany, October 5-9, 2015, Proceedings, Part III 18*, 2015, pp. 234–241.

[77] K. Chen, X. Du, B. Zhu, Z. Ma, T. Berg-Kirkpatrick, and S. Dubnov, "HTS-AT: A hierarchical token-semantic audio transformer for sound classification and detection," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2022, pp. 646–650.

[78] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative adversarial networks," *Communications of the ACM*, vol. 63, no. 11, pp. 139–144, 2020.

[79] X. Hao, J. Wu, J. Yu, C. Xu, and K. C. Tan, "Typing to listen at the cocktail party: Text-guided target speaker extraction," *arXiv preprint arXiv:2310.07284*, 2023.